

# Operability in Go

Improving operations in Go programs

# Who Am I

Ian N Schenck

SWE, Infrastructure, Oscar Health

@ianschenck (I don't really tweet)

[ian.schenck@gmail.com](mailto:ian.schenck@gmail.com)

# Preface

I am: A SWE who keeps ending up in SRE.

I try to write *operable* code.

This is going to break, how do I make it *as easy as possible*?

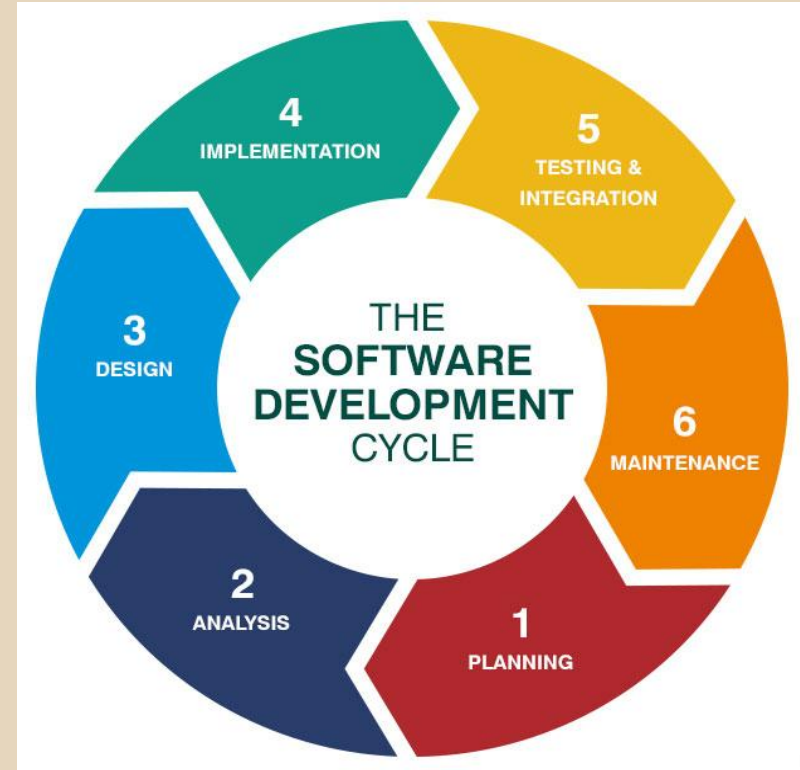
# What are Operations



# Software Design Life Cycle (SDLC)

1. Plan
2. Analyze
3. Design
4. Implement
5. Document/Test
6. Maintain

\*Not necessarily in that order



# Maintenance

Even if your software is *perfect* and *entirely bug free*, it can still break.

- Environments are complex and changing
- Hardware can break
- Humans are buggy

# Maintenance - Failure

When something fails, we have two *equal* objectives:

1. Fix it
2. Determine what went wrong

# Maintenance - Failure

## 1. Fix it

Depends on the situation.

Let's talk about failing well...



# Failing Well

- Fail immediately when unrecoverable errors occur.
- Fail the smallest execution unit necessary.
- Err on the side of caution - fail as big as you need to (maybe the whole application).

# Failing Well

In general, an unhandled/unrecoverable error should panic.

It should also give clear and concise information about what led to the panic.

# Failing Well - Panic

Applications may panic, which will fail up to a deferred `recover()`

E.g. Panic in an HTTP handler will fail up to the serving goroutine.

Panic without `recover()` terminates.

# Failing Well

Panic does give a stack trace, but could use more context.

Add context around application panics using logging.

# Maintenance - Failure

2. Determine what went wrong.

If you're unable to determine what went wrong, you can't avoid repeating the failure.

# Diagnosing Failure - 5 Whys

The vehicle will not start. (the problem)

1. Why? - The battery is dead.
2. Why? - The alternator is not functioning.
3. Why? - The alternator belt has broken.
4. Why? - The alternator belt wore out.
5. Why? - The vehicle was not maintained.  
(root cause)

# Diagnosing Failure

We need (a lot of) information!

# Killing a Stuck Process

SIGQUIT (kill -3) a process, get a stack trace:

```
goroutine 1 [IO wait, 5 minutes]:
```

```
...
```

```
net.(*TCPListener).AcceptTCP(0xc820124170, 0xc82005dbe0,  
0x0, 0x0)
```

```
    /usr/local/go/src/net/tcpsock_posix.go:254 +0x4d
```

```
net.(*TCPListener).Accept(0xc820124170, 0x0, 0x0, 0x0,  
0x0)
```

```
...
```



# Sources of Information

Stack Trace

Logs

...

# Logging

Provide context, don't just:

```
log.Println(err)
```

E.g.

```
unexpected EOF
```

Unexpected EOF of *what!*?

# A Note on Errors

Some errors provide context:

```
listen tcp :33712: bind: address already in use
```

“Named” errors (`io.ErrUnexpectedEOF`) do not:

```
unexpected EOF
```

# A Note on Errors

<https://github.com/pkg/errors>

You can add context to the error with  
`errors.Wrap`.

You can add context with the logger.

# Logging Context

Structured logging adds key-value pairs to your log.

<https://github.com/sirupsen/logrus>

```
log.WithFields(logrus.Fields{
    "animal": "walrus",
    "number": 8,
}).Debug("Started observing beach")
```

# Logging Context

Structured logging provides a way to add context in a machine and human consumable format.

```
INFO[0000] A group of walrus emerges from the ocean    animal=walrus size=10
WARN[0000] The group's number increased tremendously!  number=122 omg=true
INFO[0000] A giant walrus appears!                    animal=walrus size=10
INFO[0000] Tremendously sized cow enters the ocean.   animal=walrus size=9
FATA[0000] The ice breaks!                            number=100 omg=true
exit status 1
```

# Logging Context

Structured loggers can output text or JSON format for easy consumption by logstash/ELK/Splunk.

Context can make all the difference...

# Logging Anxiety

The anxiety over what to log, when. How much is too much, how much is enough?

Let's set this aside for now.



# Information

Other information?

Logs (action with context)

Environment

Flags

Stack Trace

More?

# Information

Logging some of these may work, but perhaps there's a better way.

Logging doesn't work at all for some cases.  
E.g. what's the current stack look like?

# Information

What about exposing information outside of logging? Logging describes *action* with *context*.

`expvar` - in the standard library. Exposes *current state*.

# expvar

Adds a route to the default ServerMux at `/debug/vars` as a side effect.

Also exposes a handler.

# expvar

`expvar` provides an http handler/endpoint which exposes arbitrary data in JSON format.

What kind of data?

- cmdline
- memstats
- And more...

# expvar

```
{  
  "cmdline": [  
    ".\./expvar_example"  
  ],  
  "memstats": {  
    "Alloc": 136736,  
    "TotalAlloc": 136736,  
    ...  
  }  
}
```

# MemStats

<https://golang.org/pkg/runtime/#MemStats>

Gives you various stats like:

- Allocated bytes (Heap/Sys/Total)
- GC statistics
- Allocations by size

# expvar

Expose various `Var` types:

- `Float*`
- `Int*`
- `Map*`
- `String`

\* Atomic operations



# expvar

Can expose a variety of “Vars”, but notably there is `Publish(Func)`:

```
func init() {  
    http.HandleFunc("/debug/vars", expvarHandler)  
    Publish("cmdline", Func(cmdline))  
    Publish("memstats", Func(memstats))  
}
```

# expvar

`expvar.Func`

Publish a function that returns `interface{}`.  
The returned value is marshalled to JSON.

# expvar

```
func memstats() interface{} {  
    stats := new(runtime.MemStats)  
    runtime.ReadMemStats(stats)  
    return *stats  
}
```

# expvar

So what can we do with `expvar`?

Expose the environment:

```
expvar.Publish("env", expvar.Func(func () interface{}  
{return os.Environ()}))
```

# Exposing Environment

## Want a map instead of a slice?

```
func publishEnv() interface{} {  
    env := make(map[string]string)  
    for _, line := range os.Environ() {  
        parts := strings.SplitN(line, "=", 2)  
        env[parts[0]] = parts[1]  
    }  
    return redactMap(env)  
}
```

# Exposing Secrets

You want to filter secret values somehow.

Replace the value with a hash so you can compare.

# Exposing Flags

Flag values can be very useful for debugging

```
func publishFlags() interface{} {  
    flagMap := make(map[string]interface{})  
    flag.VisitAll(func(f *flag.Flag) {  
        flagMap[f.Name] = f.Value  
    })  
    return redactMap(flagMap)  
}
```

# Exposing Flags

Flag magic values, for two reasons:

1. They can be changed.
2. They show up in `expvar`. (super useful)



# Exposing a Stack Trace

Publish the stack trace:

```
func publishStack() interface{} {  
    buf := make([]byte, 65535)  
    n := runtime.Stack(buf, true)  
    buf = buf[0:n]  
    return string(buf)  
}
```

# expvar ALL the things!

```
expvar.Publish("env", expvar.Func(publishEnv))  
expvar.Publish("flags", expvar.Func(publishFlags))  
expvar.Publish("stack", expvar.Func(publishStack))  
...  
expvar.Publish("my-internal-value", ...)
```



# Caution

Make sure you don't publish very expensive functions!

Don't make `expvar` too expensive or turn it into an accidental DoS vector.

# Use Verbose Names

Exposing information is great. Make that information verbose/specific:

```
expvar.Publish("jobs", ...)
```

**Better:**

```
expvar.Publish("discovery-job-cache", ...)
```

# Use Verbose Names

## Flag Names:

- `addr`
- `init-timeout`
- `rpc-retries`

## Better:

- `status-addr`, `status.addr`
- `discovery-init-timeout`, `discovery.init.timeout`
- `foo-rpc-retries`, `foo.rpc.retries`

# Use Verbose Names

## Environment Variables:

- DB\_PASSWORD
- INIT\_TIMEOUT

## Better:

- CLAIMS\_DB\_PASSWORD
- DISCOVERY\_INIT\_TIMEOUT

# expvar + structlog

These two work together.

Use expvar for *state*.

Use logging for *action*.

Less anxiety.

# More with HTTP

`expvar` means we've already committed to having a port open responding to http requests.

What else can we do with this?



# More with HTTP

- Health handler
- Specialized handlers for libraries (e.g. Vault integration)
- Shutdown handlers (quit and abort)
- Admin handlers

# Specialized Endpoints

Probably only want to do  
modification/destruction on `POST`.

`GET` can return a form.

# Specialized Endpoints - Monitoring

Prometheus is fantastic (you should use it)

Since you already have an internal/status http endpoint, dangle your prometheus metrics off of it.

# Library Developers

Provide exported variables for application developers to expose in `expvar` or logs.

Or use `expvar` and `prometheus` directly? (But that side effect)

# Recap

Think about failure at all times to guide:

- Panicking when necessary
- Exposing data via `expvar`
- Logging and context
- Naming (flags, environment variables, etc.)

# Thanks

Come talk to me about operations and go!

Questions?