

Creating a Custom Serialization Format

Scott Mansfield (@sgmansfield)
Senior Software Engineer
Netflix

What are we doing here?

1. Motivations
2. Queries
3. The Format
4. Performance
5. Future

1. Motivations

"The field is too in love with horribly inefficient frameworks. Writing network code and protocols is now considered too low level for people."

- jnordwick (Hacker News)

Motivations

- Computers make meaning out of voltages
- Serialization is everywhere
 - Network protocols
 - Video encoding
 - Machine code
 - HTTP/2 headers
 - Hard drive communication
 - Video display
- Engineers should know what's inside the black box

Motivations

- JSON is the de facto serialization format
- Common pattern:
 1. Get entire document
 2. Inflate serialized data
 3. Walk data structure & extract
- New pattern:
 1. Query the document
 2. Get only the data you need
 3. Still need to inflate

Motivations

- Query capabilities over JSON documents
- Documents stored as a byte array

JSON Document (Augmented)

```
{  
  "null"      : null,  
  "boolean"   : true,  
  "integer"   : 1,  
  "float"     : 2.3,  
  "string"    : "a string",  
  "array"     : [4, 5, 6],  
  "map"       : {"foo": 1}  
}
```


2. Queries

Query Types

- Array Index
- Array Slice
- Array Iteration
- Map Access
- Map Keys
- Map Iteration

Array Index

Query: [2]

Result: 3

[1, 2, 3, 4, 5]



Index 2

Array Slice

Query: [2:-1]

Result: [3, 4]

[1, 2, (3, 4), 5]



Index 2 until 4

Array Iteration

Query: `.a[] [0]`

Result: `[1,2,3,4,5]`

`[[1], [2], [3], [4], [5]]`

↑ ↑ ↑ ↑ ↑

Index 0 of each list

Map Access (Single)

Query: .foo

Result: 3

{"foo": 3, "bar": 4}



Key foo

Map Access (Multiple)

Query: `.foo|bar`

Result: `{"foo":3, "bar":4}`

`{"foo":3, "bar":4, "baz":5}`



Key foo



Key bar

Map Keys

Query: keys

Result: ["foo", "bar"]

{"foo": 3, "bar": 4}



Map Keys

Map Iteration

Query: `.m[] [0]`

Result: `{"foo": 3, "bar": 4}`

`{"foo": [3], "bar": [4]}`



Index 0 of each array value

Example

```
{"foo": {"k1": [3,4]},  
  "bar": {"k1": [5,6]}}
```

Query: `.m[] .k1 [0]`

Result: `{"foo": 3, "bar": 5}`

Example

```
{"foo": {"1":1, "2":2, "3":3},  
  "bar": {"4":4, "5":5, "6":6}}
```

Query: `.m[] keys`

```
Result: {"foo": ["1", "2", "3"],  
        "bar": ["4", "5", "6"] }
```

3. The Format


Types

Augmented JSON == JSON + integers

- Scalars
 - Null
 - Boolean
 - Integer (64 bit)
 - Float (64 bit)
 - String
- Composites
 - Array
 - Map

General Format

Every record starts with a single byte for the type:



Scalars

- Null
- Boolean
- Integer (64 bit)
- Float (64 bit)
- String

Scalar: Null

Type

`null`

1 byte

Scalar: Boolean


Type	Data
bool	1 or 0
1 byte	1 byte

Scalar: Integer

Type	Data
int	Little endian int64
1 byte	8 bytes

Scalar: Integer (example)

4 = 0x0000_0000_0000_0004




Scalar: Float

Type	Data
float	float64 as little endian uint64
1 byte	8 bytes

Scalar: Float (example)

4.5 = 0x4012_0000_0000_0000




Scalar: String

Type	Length	Data
string	Little endian uint32	String contents
1 byte	4 bytes	length bytes

Scalar: String (example)

"Hello, Go!" Length: 10 = 0x0000_000A

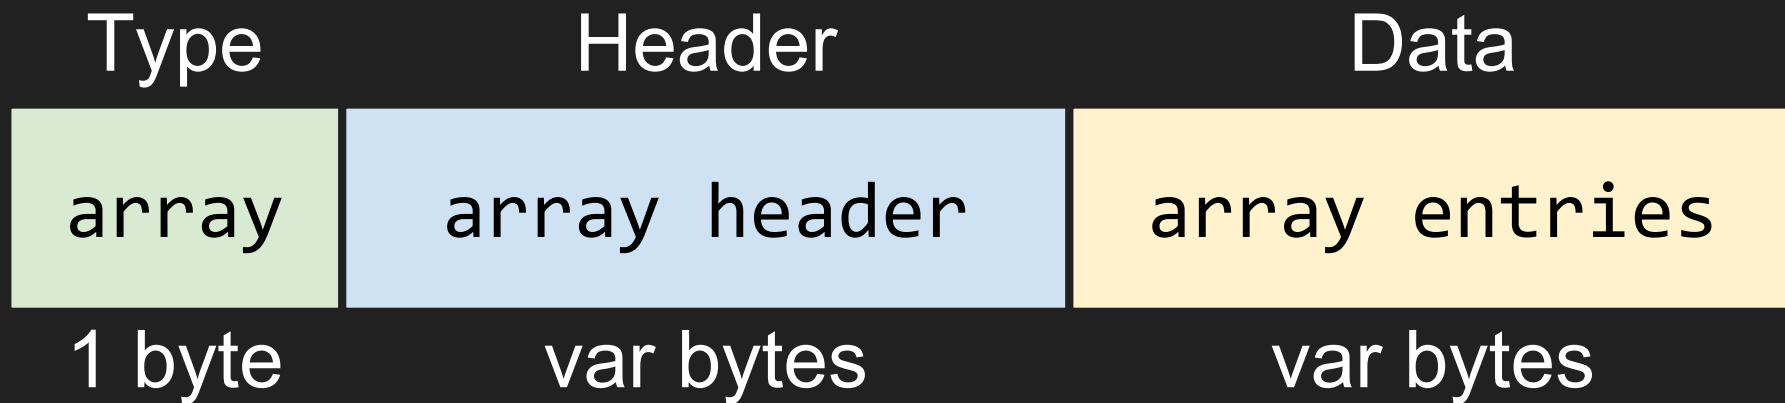


Composites

Recursive - contained data are defined by this same format

- Array
- Map

Composite: Array




Composite: Array - Header

numoffsets	offlen	offsets
uvarint	(0,8)	numoffsets uints of offlen length
var bytes	1 byte	numoffsets × offlen bytes

Composite: Array - Header offsets


2 or more offsets



Each offset is 4 bytes

Composite: Array - Data

1 or more records



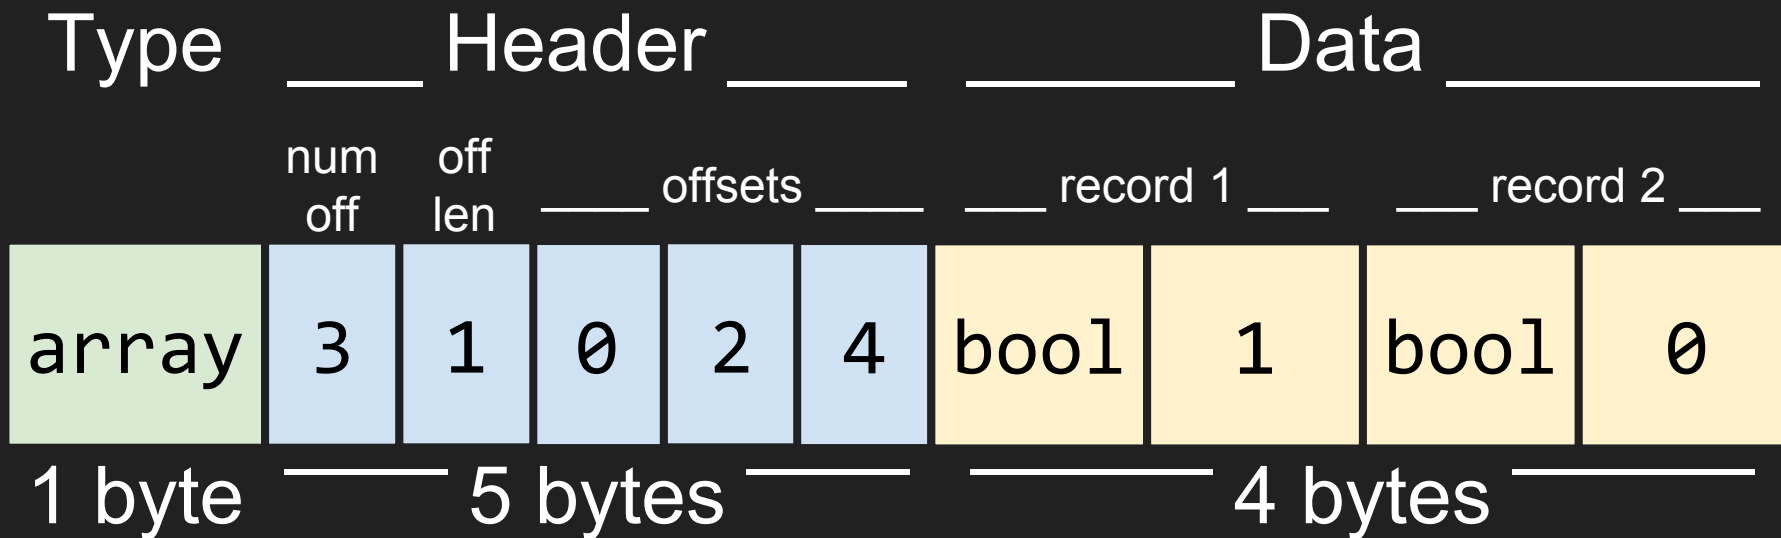
Each var bytes

Composite: Empty Array

Type	numoffsets
array	uvarint (0)
1 byte	1 byte

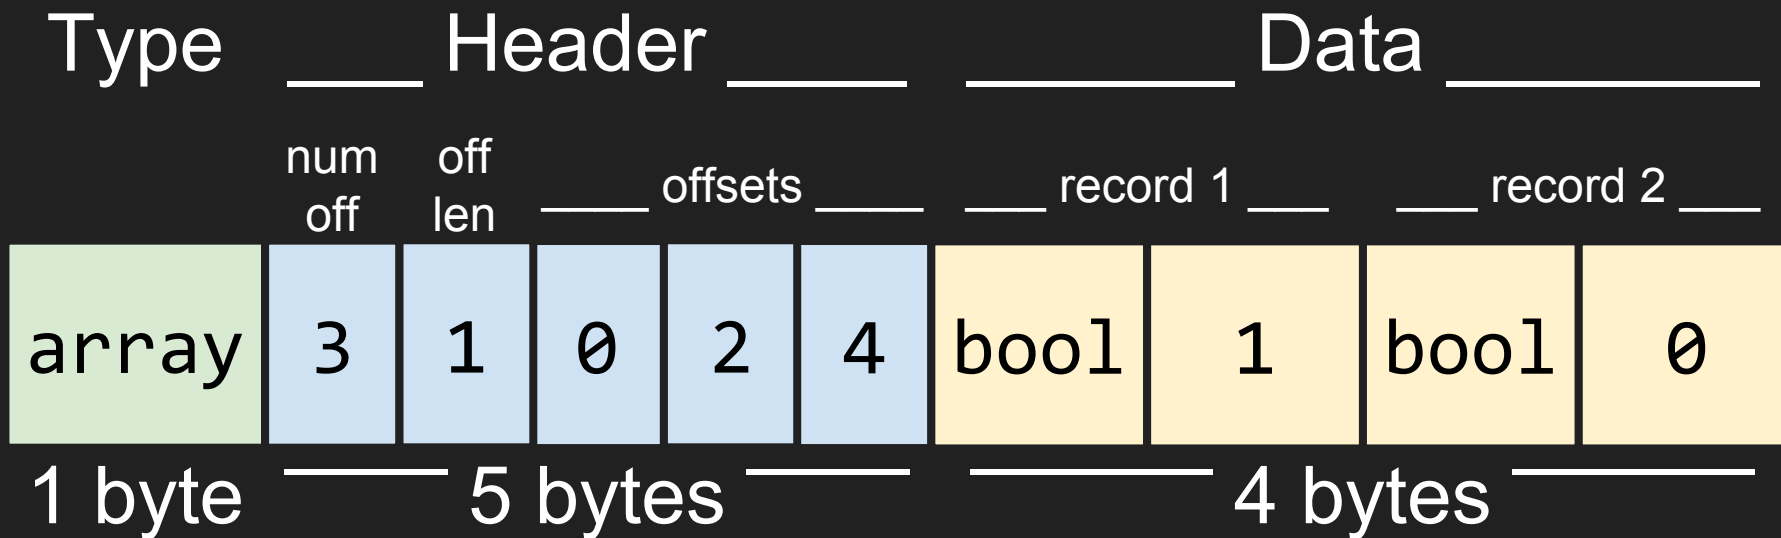
Composite: Array (example)

[true, false]




Composite: Array (example, slicing)

[true, false]



Composite: Map




Composite: Map - Header

num recs	offlen	lenlen	header records
uvarint	(0, 8)	(0, 8)	num recs header records
var bytes	1 byte	1 byte	\propto num recs

Composite: Map - Header


1 or more header records



Each 4 + offlen + lenlen bytes

Composite: Map - Data

1 or more records



Each var bytes

Composite: Map - Header Record

Intern ID	offset	length
uint32	uint	uint
4 bytes	offset bytes	length bytes

Composite: Map - Interned Keys

- Map keys are assigned a unique uint32 ID
- IDs are shared by identical strings
- Forward and reverse mappings stored next to the data
- Example:
 - "true" → 1
 - "false" → 2

Composite: Map - Header

header records

1

17

52

195

Composite: Empty Map

Type	num recs
map	uvarint (0)
1 byte	1 byte

"false" → 2

```
{"false":false, "true":true}
```

Type	Header	Data
------	--------	------

header records

rec

off
len


len
len

record 1

record 2

record 1

record 2



4. Performance

How fast is it?

It depends

... on:

- How much data you ask for
- How complex the query is
- How many CPU's
- Speed of the underlying data storage

Scalars

Serialize

Type	time/op
Null	64.3 ns \pm 2%
Boolean	71.6 ns \pm 1%
Int	75.7 ns \pm 0%
Float	75.4 ns \pm 1%
String	88.6 ns \pm 1%

"foobar"

Deserialize

Type	time/op
Null	16.0 ns \pm 1%
Boolean	23.9 ns \pm 1%
Int	26.6 ns \pm 1%
Float	27.1 ns \pm 1%
String	70.1 ns \pm 1%

Composites: Serialize

Type	# elems	time/op	time/op (ns)
Array	0	115 ns \pm 0%	115 ns
Array	1	273 ns \pm 1%	273 ns
Array	10	900 ns \pm 1%	900 ns
Array	100	5.42 μ s \pm 1%	5420 ns
Array	1000	43.7 μ s \pm 1%	43700 ns
Array	10000	453 μ s \pm 1%	453000 ns
Array	100000	5.35 ms \pm 1%	5350000 ns
Array	1000000	54.0 ms \pm 3%	54000000 ns
Map	0	87.2 ns \pm 1%	87 ns
Map	1	608 ns \pm 1%	608 ns
Map	10	3.39 μ s \pm 1%	3390 ns
Map	100	34.1 μ s \pm 1%	34100 ns
Map	1000	374 μ s \pm 0%	374000 ns
Map	10000	4.37 ms \pm 1%	4370000 ns
Map	100000	58.7 ms \pm 2%	58700000 ns
Map	1000000	866 ms \pm 4%	866000000 ns

Composites: Deserialize

Type	# elems	time/op	time/op (ns)
Array	0	136 ns \pm 1%	136 ns
Array	1	201 ns \pm 0%	201 ns
Array	10	588 ns \pm 2%	588 ns
Array	100	4.05 μ s \pm 3%	4050 ns
Array	1000	38.1 μ s \pm 1%	38100 ns
Array	10000	380 μ s \pm 2%	380000 ns
Array	100000	3.81 ms \pm 1%	3810000 ns
Array	1000000	39.9 ms \pm 2%	39900000 ns
Map	0	158 ns \pm 0%	158 ns
Map	1	361 ns \pm 0%	361 ns
Map	10	1.97 μ s \pm 0%	1970 ns
Map	100	21.3 μ s \pm 0%	21300 ns
Map	1000	261 μ s \pm 1%	261000 ns
Map	10000	2.67 ms \pm 1%	2670000 ns
Map	100000	38.3 ms \pm 2%	38300000 ns
Map	1000000	757 ms \pm 3%	757000000 ns

Composites: Queries

Type	# elems	time/op
Array Get	1	25.9 ns \pm 7%
Array Get	10	26.4 ns \pm 6%
Array Get	100	26.6 ns \pm 6%
Array Get	1000	26.3 ns \pm 6%
Array Get	10000	26.3 ns \pm 8%
Array Get	100000	26.0 ns \pm 4%
Array Get	1000000	26.2 ns \pm 7%

Map Get	1	35.3 ns \pm 1%
Map Get	10	64.7 ns \pm 0%
Map Get	100	74.6 ns \pm 1%
Map Get	1000	121 ns \pm 1%
Map Get	10000	157 ns \pm 0%
Map Get	100000	221 ns \pm 2%
Map Get	1000000	375 ns \pm 1%

Type	# elems	time/op
Array Slice	1	70.1 ns \pm 1%
Array Slice	10	73.9 ns \pm 4%
Array Slice	100	73.7 ns \pm 3%
Array Slice	1000	73.0 ns \pm 2%
Array Slice	10000	73.4 ns \pm 3%
Array Slice	100000	75.6 ns \pm 3%
Array Slice	1000000	73.4 ns \pm 2%

Map Keys	1	662 ns \pm 9%
Map Keys	10	2.11 μ s \pm 8%
Map Keys	100	17.4 μ s \pm 8%
Map Keys	1000	173 μ s \pm 8%
Map Keys	10000	2.28 ms \pm 4%
Map Keys	100000	35.6 ms \pm 5%
Map Keys	1000000	348 ms \pm 7%

5. Future

In Progress & Future Work

- Replace simple scalar values
- Append to arrays
- Add new keys to a map
- Other ops (inc, dec, etc)
- Compression

Thank You

@sgmansfield
smansfield@netflix.com
techblog.netflix.com

