

Go at the DARPA Cyber Grand Challenge

Channels and Parallelism for High Performance
Database, Network and File I/O

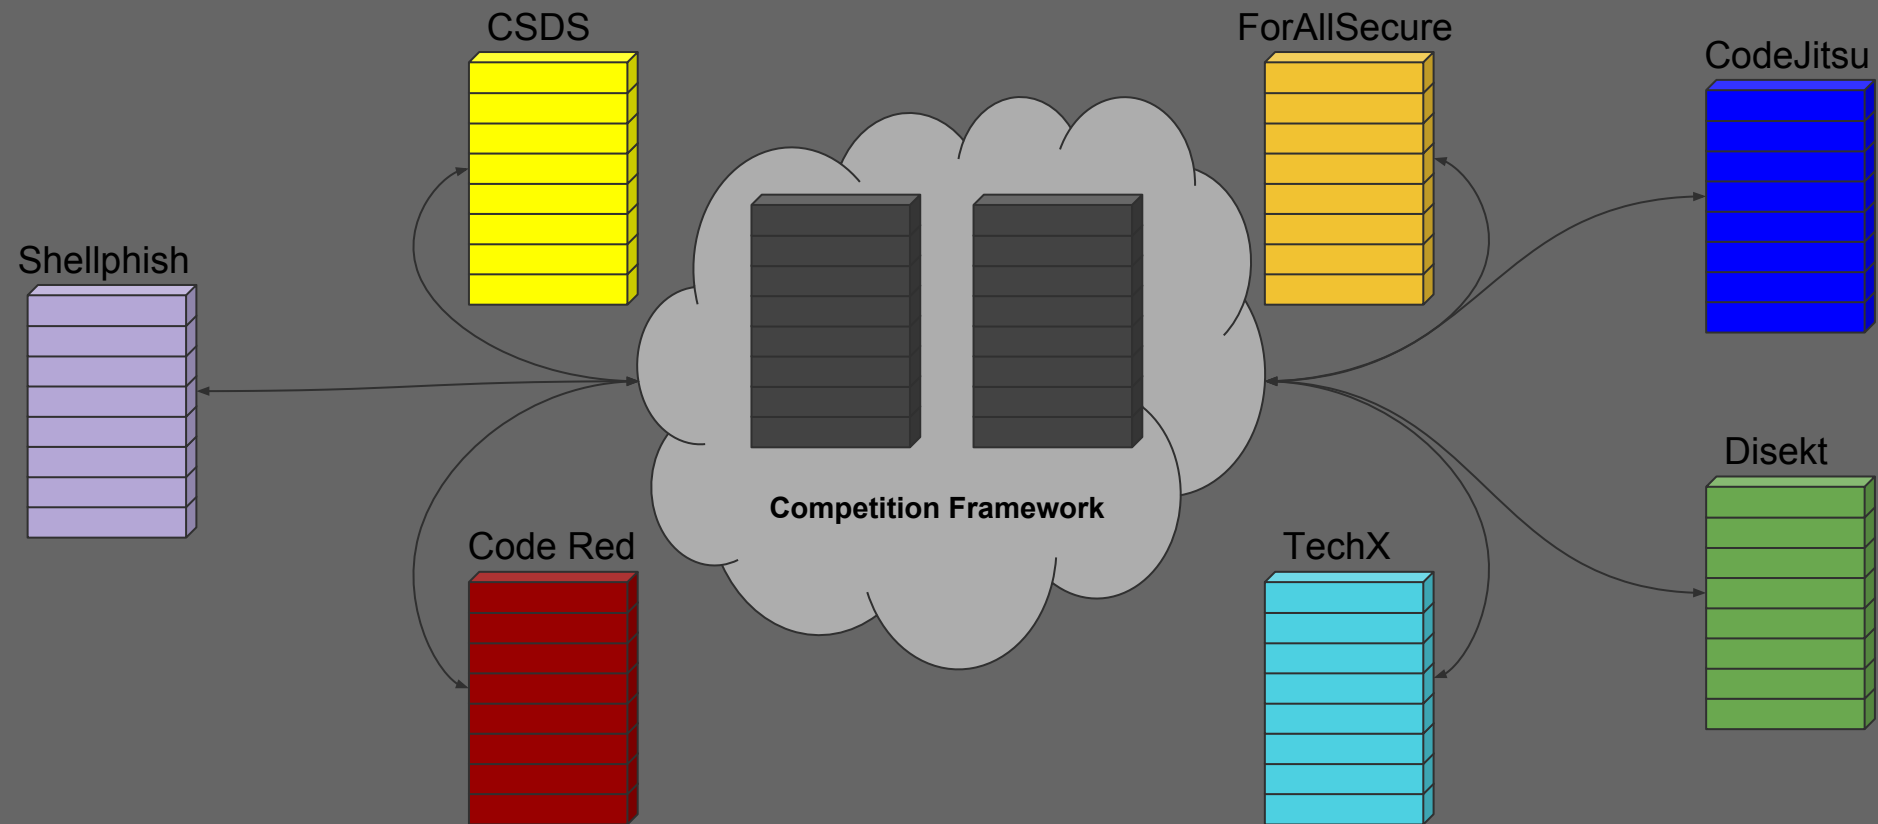
Will Hawkins

Introduction

DARPA Cyber Grand Challenge



DARPA Cyber Grand Challenge



DARPA Cyber Grand Challenge




Our Goal: Finish 1000

DARPA Cyber Grand Challenge



Our Goal: Finish 1000 (= 0111 = “Not 7th”)

DARPA Cyber Grand Challenge



Goal

- Get it? The go-al?
- Capturer will capture ...
 - ... the fields from the packets ...
 - Timestamp
 - Cbid
 - Conversation
 - Side - either client or server
 - Message ID
 - Contents
 - ... into a data store that can be queried by different components.
 - The Fuzzer
 - The Rule Generator

Results

- This ended up being a significant component of our defensive systems
- The Fuzzer
 - Used data captured from the network tap to generate inputs for our Fuzzers
- The Rule Generator
 - Rules were deployed in concert with hardened binaries.
 - Hard to tell which was the effective defense.
 - However, there were two specific cases where replacement binaries were vulnerable *but* the IDS rules generated by the Rule Generator protected the binary from successful attack.

Design, Architecture and Implementation

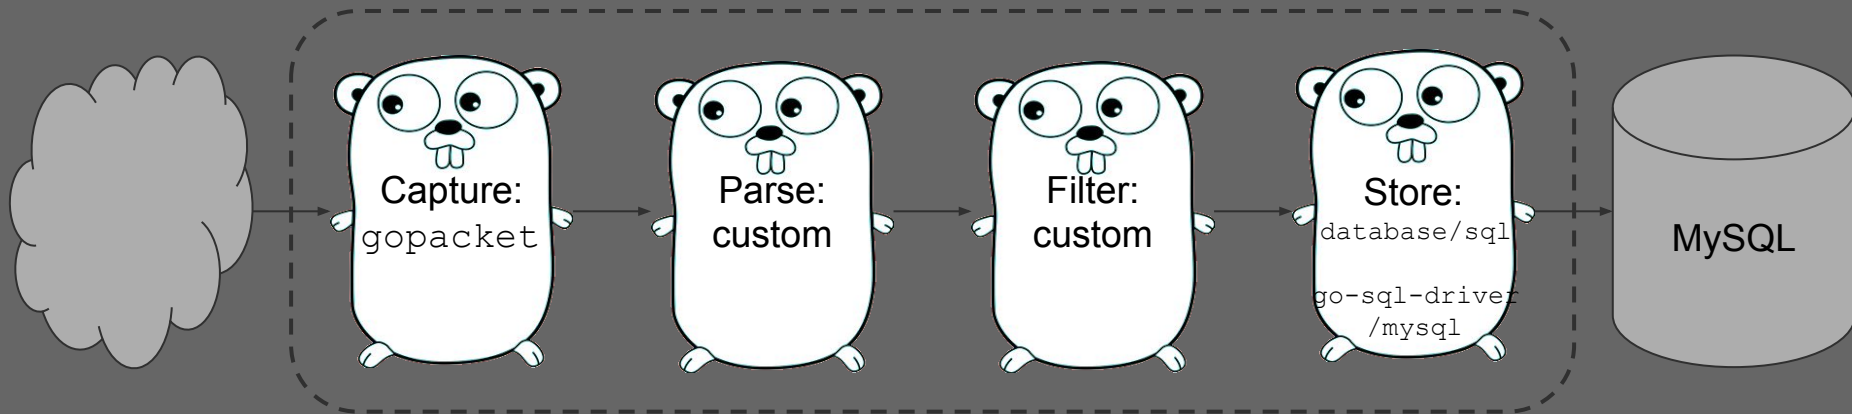
- Capturing
 - Simple tcpdump
 - libpcap
- Filtering
 - BPF
 - Custom
- Storing
 - Flat file
 - Database
 - SQL
 - NoSQL
- Off-the-shelf? Custom?
 - BASH pipeline: `tcpdump 'port 1993' | sed | awk | ... > file`
 - Custom program
 - C
 - C++
 - Go

Design, Architecture and Implementation

- Capturing
 - Simple tcpdump
 - **libpcap**
- Filtering
 - BPF
 - **Custom**
- Storing
 - Flat file
 - **Database**
 - **SQL: MySQL**
 - **NoSQL**
- Off-the-shelf? Custom?
 - BASH pipeline: tcpdump 'port 1993' | sed | awk | ... > file
 - Custom program
 - **C**
 - **C++**
 - **Go**

Design, Architecture and Implementation

- Capturing
- Filtering
- Storing
- Off-the-shelf? Custom?



Capturing

- gopacket
 - API for
 - accessing different packet sources in a consistent manner (`PacketDataSource` interface), and
 - accessing data in packets in a consistent manner
 - Sources include: byte array, files (from, e.g., `tcpdump` output), live devices (from, e.g., `libpcap`, `pfring`)
- Three steps for use:
 - Handle
 - `OpenLive()`
 - `OpenOffline()`
 - Source from Handle
 - `NewPacketSource()`
 - Capture
 - `Packets()`

Capturing

```
live_handle, err = pcap.OpenLive(*iface, 1024, false, pcap.BlockForever)
...
defer live_handle.Close()
...
live_source := gopacket.NewPacketSource(live_handle, live_handle.LinkType())
...
for stopping == false {
    select {
        ...
        case p, ok := <- live_source.Packets():
            if !ok {
                stopping = true
                break
            }
    }
}
```

Filtering

- gopacket (again)
- Two Steps For Use:
 - From a packet, get the Layer: `Layer()`
 - Use the Layer

Filtering

```
var ipv4_layer gopacket.Layer
var udp_layer gopacket.Layer
var udp *layers.UDP
var ip *layers.IPv4

if ipv4_layer = i.Layer(layers.LayerTypeIPv4); ipv4_layer == nil {
    ...
}
if udp_layer = i.Layer(layers.LayerTypeUDP); udp_layer == nil {
    ...
}
udp, _ = udp_layer.(*layers.UDP)
ip, _ = ipv4_layer.(*layers.IPv4)
if filter.DstPorts != nil && !cgc_utils.MatchAnyPort(udp.DstPort, filter.DstPorts) {
    ...
}
```


Parsing

```
func ParseCgcPacket(packet []byte) (cgc_packet IdsPacket, err error) {
    packet_length := len(packet)
    packet_offset := 0

    if (packet_offset+4) > packet_length {
        err = errors.New("Could not parse past first field.")
        return
    }
    csid := binary.LittleEndian.Uint32(packet[packet_offset:])
    cgc_packet.Csid = fmt.Sprintf("%x", csid)
    packet_offset += 4

    if (packet_offset+4) > packet_length {
        err = errors.New("Could not parse past csid field.")
        return
    }
    cgc_packet.ConnectionID = binary.LittleEndian.Uint32(packet[packet_offset:])
    packet_offset += 4
    ...
    return
}
```

Storing

- database/sql in combination with Go-MySQL-Driver
- Three steps for use:
 - Connect to the database: `Open()`
 - Prepare statements (optional): `Prepare()`
 - Execute statements: `Exec()`
 - Check/Retrieve Results:
 - Check return value, or
 - `Next()`

Storing

```
db, err := sql.Open("mysql",
                    user+":"+password+"@unix(/var/run/mysqld/mysqld.sock)/"+database)

if err != nil {
    return nil, err
}

if err = db.Ping(); err != nil {
    return nil, err
}

statement, statement_err = database.Prepare("insert into pcap
                                           (cbid, conversation, side, message, contents)
                                           VALUES (?, ?, ?, ?, ?)")

if _, err := statement.Exec(packet.Csid,
                             packet.ConnectionID,
                             fmt.Sprintf("%v", packet.ConnectionSide),
                             packet.MessageID,
                             packet.Message); err != nil {
    fmt.Printf("Exec() error: %v\n", err)
}
```

Testing

- DARPA played games with us
- `tcpdump`: Capture the traffic from the simulation
 - All
 - Some
- `tcpreplay`
 - Real time
 - Accelerated
- Storage required
 - Fixed
 - Variable

Testing

- DARPA played games with us
- `tcpdump`: Capture the traffic from the simulation
 - All
 - **Some**: Filter meta traffic (ssh, etc); Use only 200,000 packets.
- `tcpreplay`
 - Real time
 - **Accelerated**: `-t`; Playback all 200,000 packets in 15s; That's roughly 7.43Mbps.
- Storage required
 - Fixed
 - Variable
 - **N/A**

Initial Result

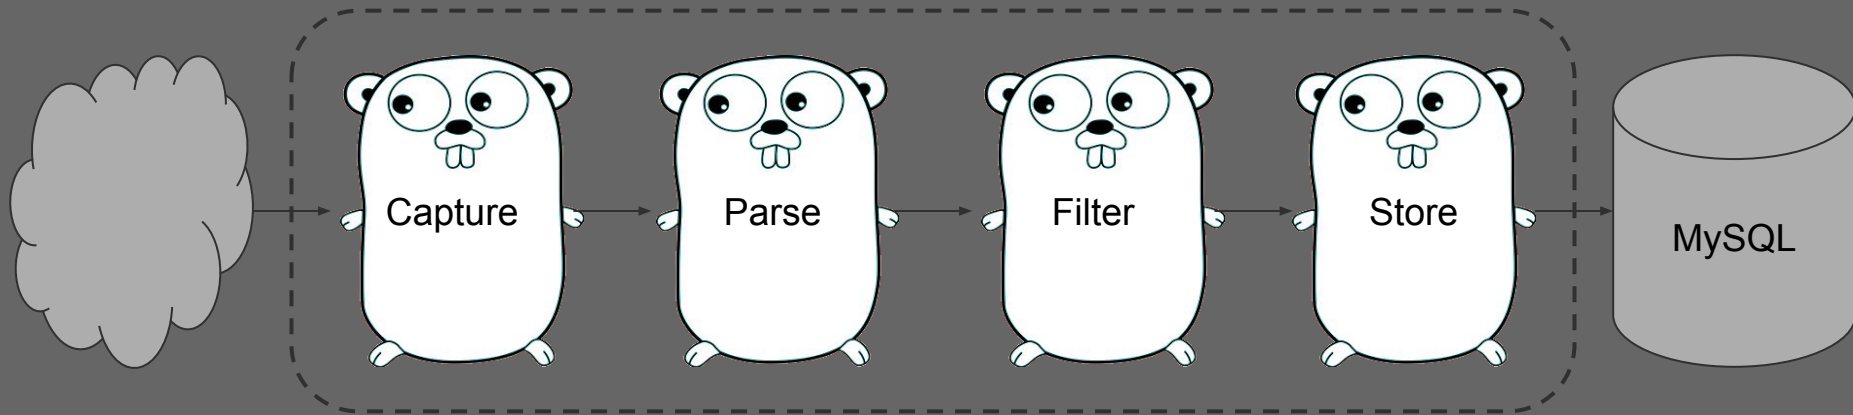
$200,000\text{pkts}/200,000\text{pkts}=\mathbf{100\%}$

Questions


Initial Result

$$844\text{pkts}/200,000\text{pkts}=\mathbf{0.42\%}$$


Optimization




Optimization



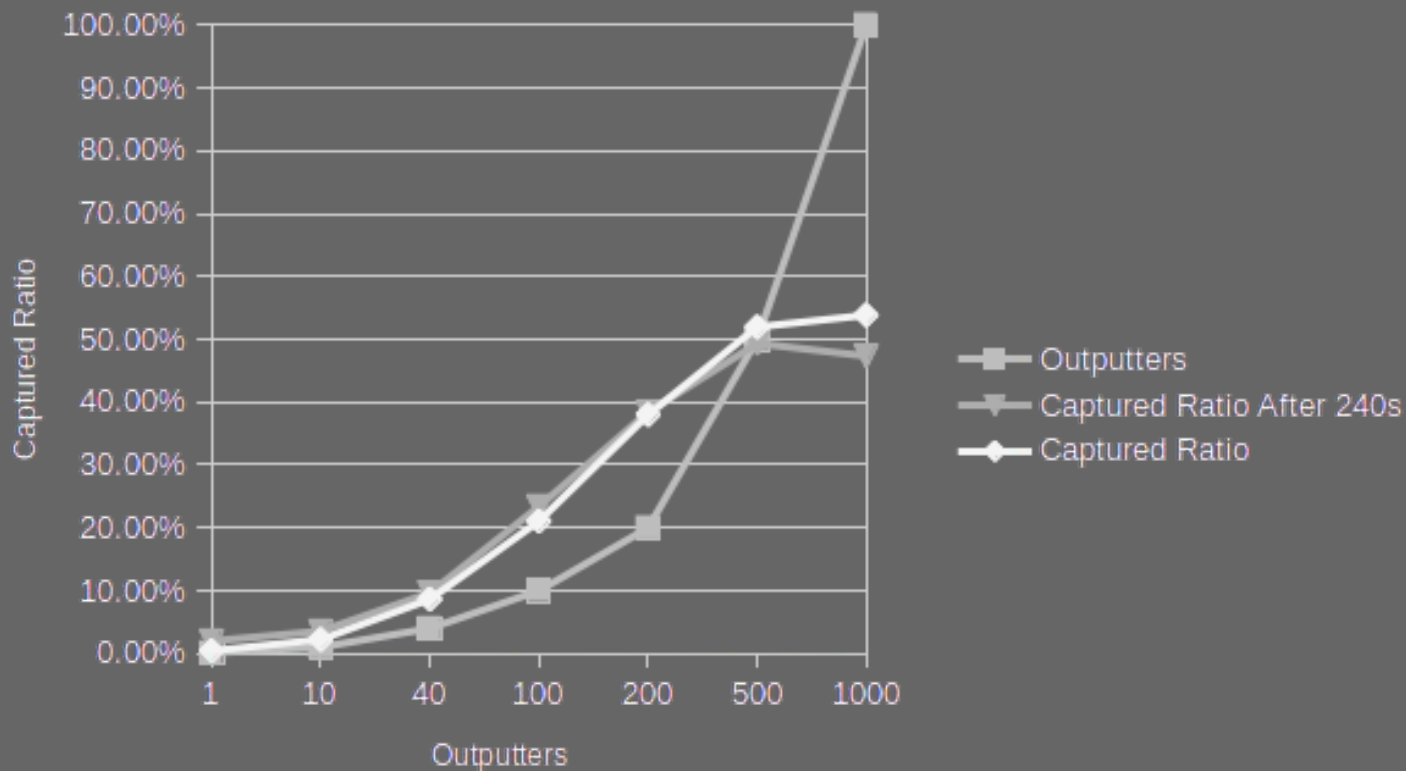
Optimization



Outputters' Effect on Capture Ratios



Outputters' Effect on Capture Ratios Over Time



Buffer(s)

- There appears to be buffering
- Sources
 - Capturer itself
 - `gopacket / libpcap`
 - Operating system

Buffer(s)

- There appears to be buffering
- Sources
 - **Capturer itself: Nope**
 - `gopacket / libpcap`
 - Operating system

Buffer(s)

- There appears to be buffering
- Sources
 - Capturer itself
 - **gopacket / libpcap**

```
func (p *PacketSource) Packets() chan Packet {  
    if p.c == nil {  
        p.c = make(chan Packet, 1000)  
        go p.packetsToChannel()  
    }  
    return p.c  
}
```

- Operating system

Buffer(s)

- There appears to be buffering
- Sources
 - Capturer itself
 - **gopacket / libpcap**

```
func (p *PacketSource) Packets() chan Packet {  
    if p.c == nil {  
        p.c = make(chan Packet, 1000)  
        go p.packetstoChannel()  
    }  
    return p.c  
}
```

- Operating system

Buffer(s)

- There appears to be buffering
- Sources
 - Capturer itself
 - `gopacket / libpcap`
 - **Operating system**

On Linux, `rmem_default`:

```
$ cat /proc/sys/net/core/rmem_default  
212992
```

On BSD, `net.bpf.size/maxsize`.

Optimization

Profiling

- Misleading
- Helpful
- In Between

Profiling

- Misleading

```
go tool pprof -cum -top cap pcap.prof | head
```

```
210ms of 210ms total ( 100%)
```

flat	flat%	sum%	cum	cum%	
0	0%	0%	140ms	66.67%	runtime.goexit
0	0%	0%	120ms	57.14%	github.com/google/gopacket .(*PacketSource).NextPacket
0	0%	0%	120ms	57.14%	github.com/google/gopacket .(*PacketSource).packetsToChannel
0	0%	0%	110ms	52.38%	github.com/google/gopacket /pcap.(*Handle).ReadPacketData
0	0%	0%	110ms	52.38%	github.com/google/gopacket /pcap.(*Handle).getNextBufPtrLocked
10ms	4.76%	4.76%	110ms	52.38%	github.com/google/gopacket /pcap._Cfunc_pcap_next_ex
80ms	38.10%	42.86%	100ms	47.62%	runtime.cgocall
0	0%	42.86%	40ms	19.05%	runtime._System

- Helpful
- In Between

Profiling

- Misleading
- Helpful

```
$ go tool pprof -cum -top cap pcap.prof | head
```

```
80ms of 80ms total ( 100%)
```

	flat	flat%	sum%		cum	cum%	
	0	0%	0%		50ms	62.50%	runtime.goexit
	0	0%	0%		30ms	37.50%	runtime.gcDrain
	0	0%	0%		20ms	25.00%	runtime.gcBgMarkWorker
	20ms	25.00%	25.00%		20ms	25.00%	runtime.scanobject
	10ms	12.50%	37.50%		20ms	25.00%	runtime.systemstack
	0	0%	37.50%	10ms	12.50%		database/sql.(*Stmt).Exec
	0	0%	37.50%	10ms	12.50%		database/sql.resultFromStatement
	0	0%	37.50%	10ms	12.50%		github.com/go-sql-driver/mysql.(*buffer).fill

- In Between

Profiling

- Misleading
- Helpful
- In Between

```
$ go tool pprof -cum -top cap pcap.prof | head
```


```
730ms of 1120ms total (65.18%)
```

```
Showing top 80 nodes out of 150 (cum >= 20ms)
```

	flat	flat%	sum%	cum	cum%	
0	0%	0%	810ms	72.32%	runtime.goexit	
0	0%	0%	360ms	32.14%	github.com/google/gopacket.(*PacketSource).NextPacket	
0	0%	0%	360ms	32.14%	github.com/google/gopacket.(*PacketSource).packetsToChannel	
0	0%	0%	300ms	26.79%	github.com/google/gopacket/pcap.(*Handle).ReadPacketData	
0	0%	0%	300ms	26.79%	github.com/google/gopacket/pcap.(*Handle).getNextBufPtrLocked	
0	0%	0%	300ms	26.79%	main.db_extract_output	
0	0%	0%	280ms	25.00%	github.com/google/gopacket/pcap._Cfunc_pcap_next_ex	

Profiling

- What is the cause of bogus profiling data?
- Think about the order of events



Profiling

- What is the cause of bogus profiling data?
- Think about the order of events

time	Capture	Parse	Filter	Store
0	read()	select()	select()	select()
1	read()	select()	select()	select()
2	read()	select()	select()	select()
3	read()	select()	select()	select()
4	read()	select()	select()	select()



Working.

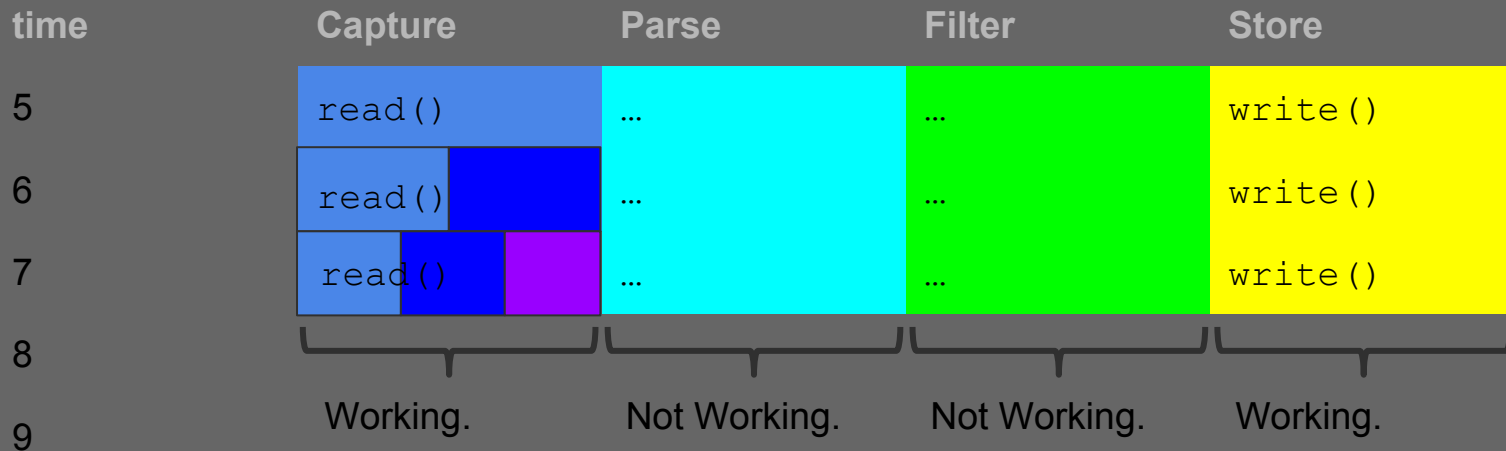
Profiling

- What is the cause of bogus profiling data?
- Think about the order of events

	time	Capture	Parse	Filter	Store
	0	read()	select()	select()	select()
Data starts →	1	read()	select()	select()	select()
	2	read()	receive()	select()	select()
	3	read()	receive()	receive()	select()
	4	read()	receive()	receive()	receive()

Profiling

- What is the cause of bogus profiling data?
- Think about the order of events



Profiling

- What is the cause of bogus profiling data?
- Think about the order of events

	time	Capture	Parse	Filter	Store
	0	read()	select()	select()	select()
Data starts →	1	read()	select()	select()	select()
	2	read()	receive()	select()	select()
	3	read()	receive()	receive()	select()
	4	read()	receive()	receive()	receive()

Manual Profiling

- The equivalent of `printf()` debugging
- Remove components until performance improves

Manual Profiling

- The equivalent of `printf()` debugging
- Remove components until performance improves

**THE BEATINGS WILL
CONTINUE UNTIL
MORALE IMPROVES.**

Manual Profiling

- The equivalent of `printf()` debugging
- Remove components until performance improves


```
func log_cgc_packet(packet cgc_utils.IdsPacket, statement sql.Stmt) {  
    _, err := statement.Exec(packet.Csid,  
                             packet.ConnectionID,  
                             fmt.Sprintf("%v", packet.ConnectionSide),  
                             packet.MessageID,  
                             packet.Message)  
  
    if err != nil {  
        fmt.Printf("Exec() error: %v\n", err);  
    }  
}
```

Manual Profiling


- The equivalent of `printf()` debugging
- Remove components until performance improves

```
func log_cgc_packet(packet cgc_utils.IdsPacket, statement sql.Stmt) {  
    return  
    if _, err := statement.Exec(packet.Csid,  
                                packet.ConnectionID,  
                                fmt.Sprintf("%v", packet.ConnectionSide),  
                                packet.MessageID,  
                                packet.Message); err != nil {  
        fmt.Printf("Exec() error: %v\n", err)  
    }  
}
```

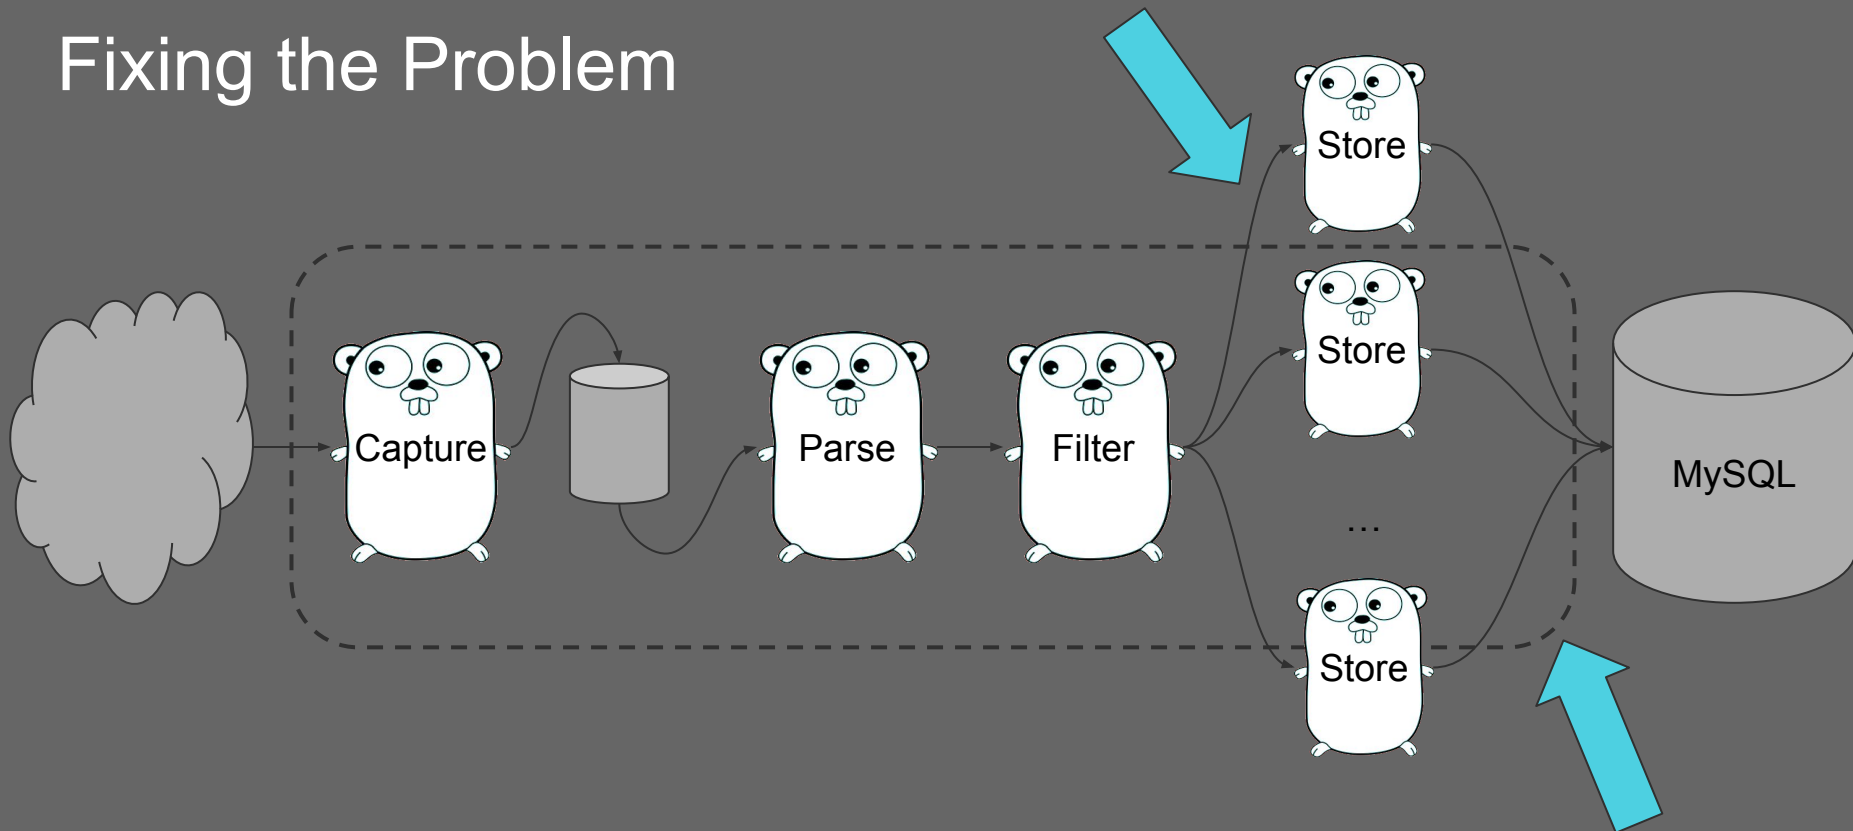

Fixing the Problem



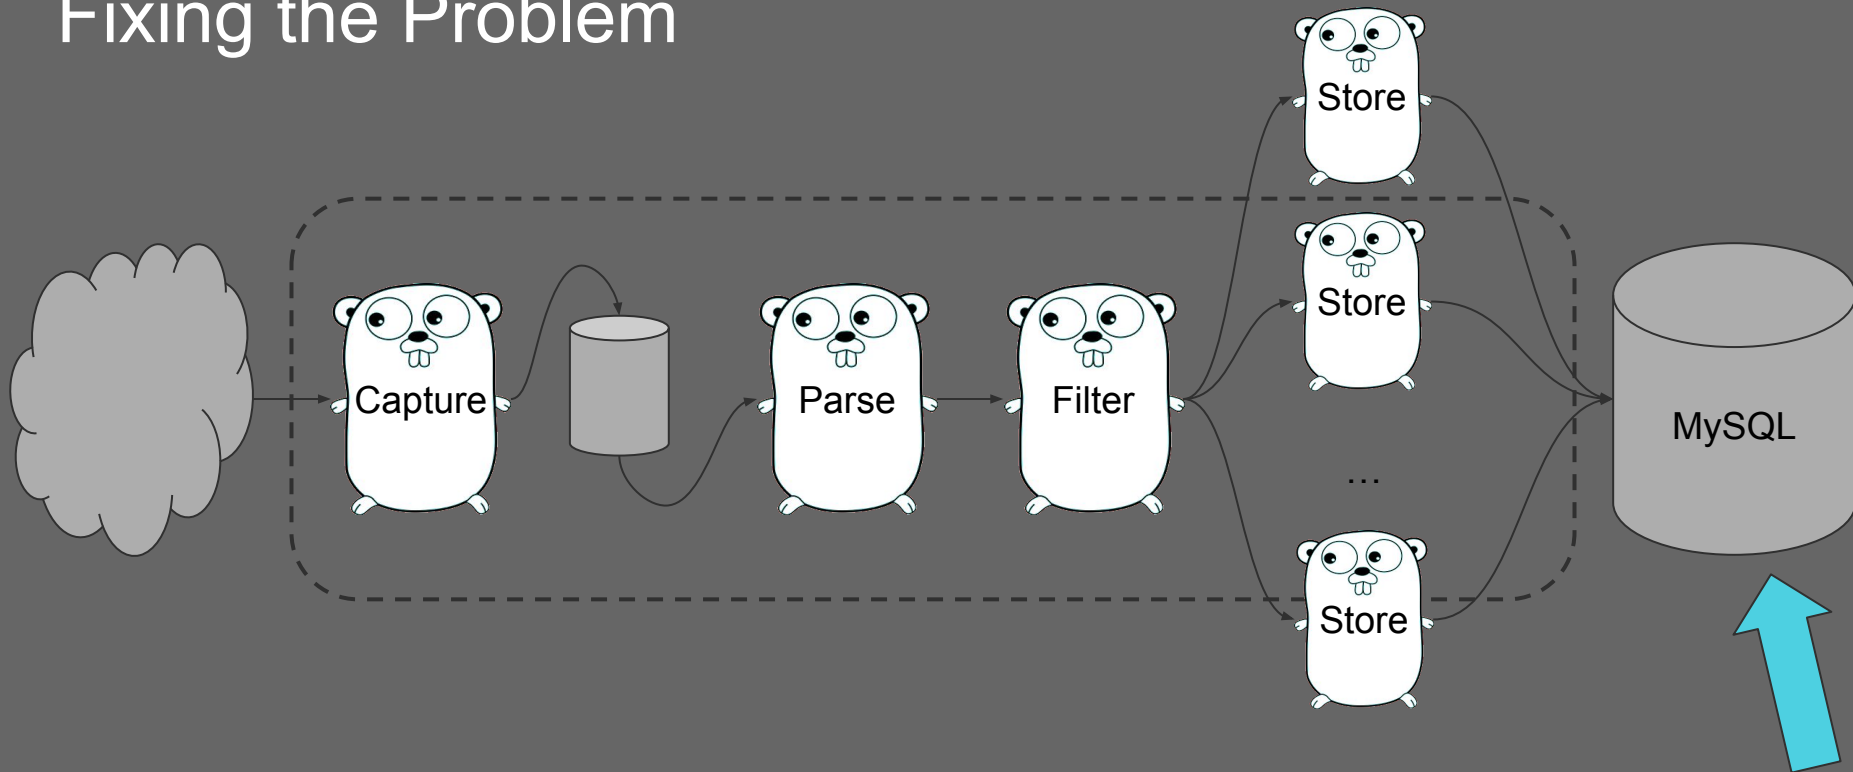
Fixing the Problem



Fixing the Problem



Fixing the Problem



#winning



Questions